



kubernetes

Why kubernetes ?



Container Clusters

- What if we have 10s, 100s, 1000s of running containers on multiple VMs?
- How to deploy, scale, restart, manage all of these containers?
- What problems do they solve?
 - Management
 - Metrics
 - Health checks
 - Security
 - Abstraction of hardware
 - Networking
 - Scheduling
 - Scaling
 - Deployment
 - Rollbacks
 - Zero-downtime / blue-green
 - Service discovery

A Brief Kubernetes History

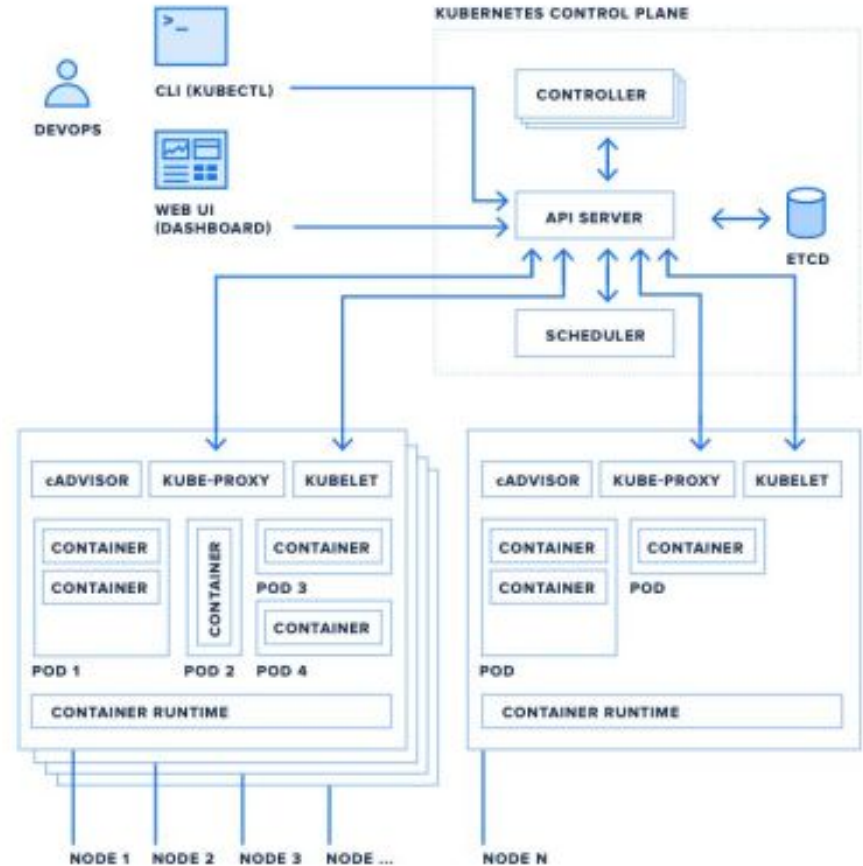
- “K8s”
- Evolved out of Borg (Google’s internal container cluster)
- Open sourced ~2014
- Grew in popularity, open source velocity increased
- Now the most popular container cluster (most cloud platforms have some sort of managed K8s offering)
- Features added regularly and frequently
- Cloud Native / CNCF - Kubernetes, Prometheus, Fluentd

Kubernetes Architecture

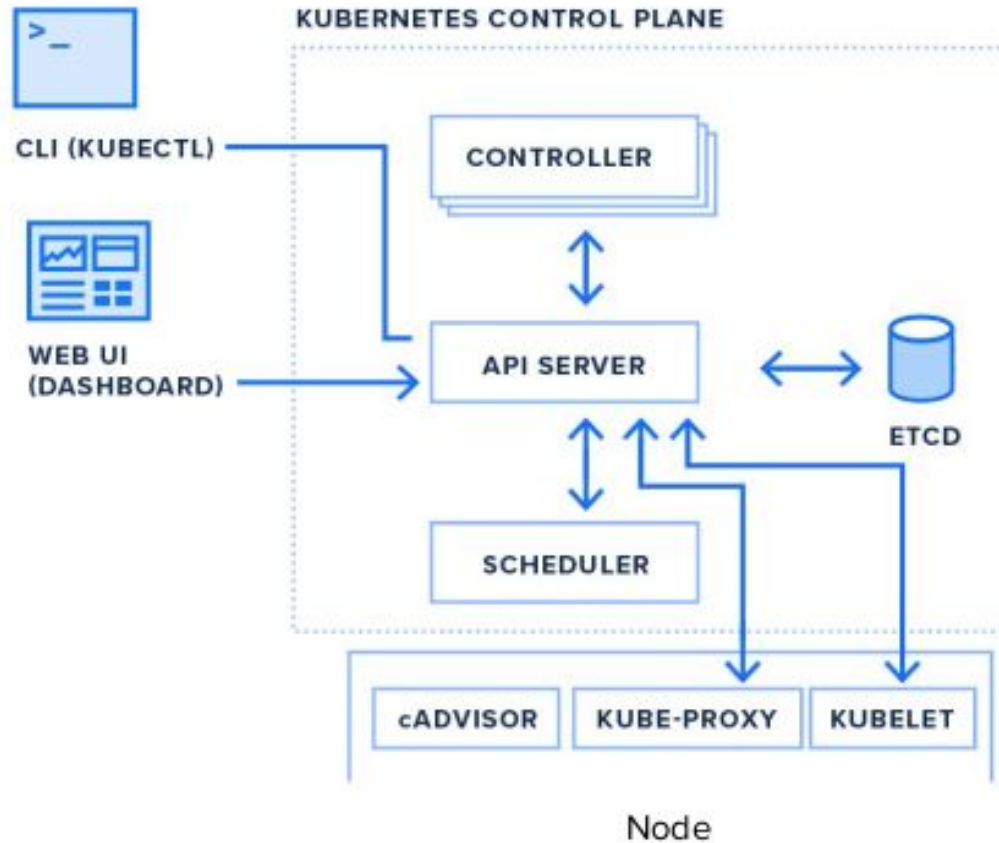
- Client side :- CLI(KUBECTL)

WEB UI(DASHBOARD)

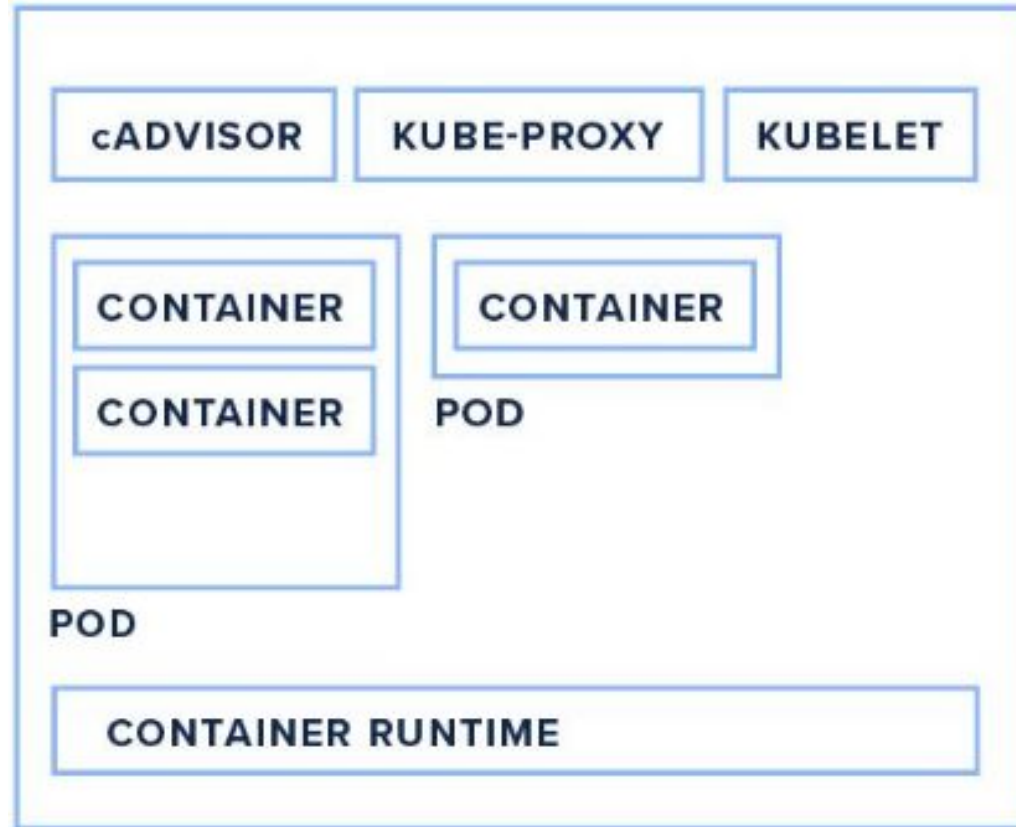
- Kubernetes Control plane or master node
- Minions node or worker node



Kubernetes control plane



Kubernetes worker nodes



Kubernetes installation

Single Node

- Docker desktop
- Minikube

Custom kubernetes

- Kubeadm
- Kubespray

Cloud

- AWS - EKS
- Azure - AKS
- Google - GKE

Minikube installation

What you'll need

- 2 CPUs or more
- 2 GB of free memory
- 20 GB of free disk space
- Internet connection
- Container or virtual machine manager, such as: [Docker](#), [QEMU](#), [Hyperkit](#), [Hyper-V](#), [KVM](#), [Parallels](#), [Podman](#), [VirtualBox](#), or [VMware Fusion/Workstation](#)

FOR linux

```
curl -LO https://storage.googleapis.com/minikube/releases/latest/minikube-linux-amd64
```

```
sudo install minikube-linux-amd64 /usr/local/bin/minikube
```

For windows

<https://storage.googleapis.com/minikube/releases/latest/minikube-installer.exe>

If using powershell

```
New-Item -Path 'c:\' -Name 'minikube' -ItemType Directory -Force
```

```
Invoke-WebRequest -OutFile 'c:\minikube\minikube.exe' -Uri  
'https://github.com/kubernetes/minikube/releases/latest/download/minikube-windows-amd64.exe'  
-UseBasicParsing
```

FOR mac

curl -LO

<https://storage.googleapis.com/minikube/releases/latest/minikube-darwin-amd64>

sudo install minikube-darwin-amd64 /usr/local/bin/minikube

Some K8s commands

Minikube start

Minikube stop

Kubectl version

Kubectl get

Kubectl apply


Kubectl create

Kubectl delete

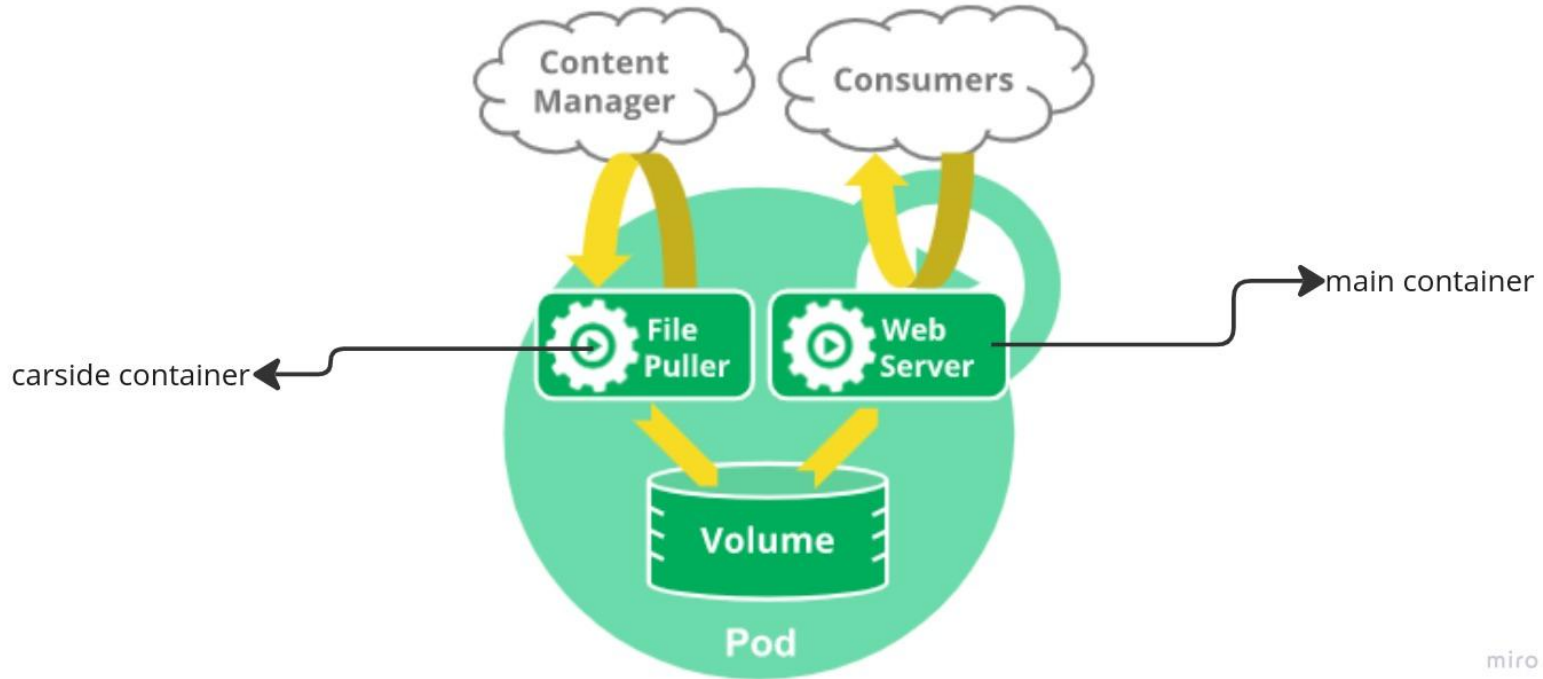
K8s components

PODS

- **Pods** are the smallest deployable units of computing that you can create and manage in Kubernetes.
- A **Pod** is a group of one or more containers.
- **Pods** that run a single container.
- **Pods** that run multiple containers that need to work together.
- **Pod** containers share resources
 - Storage
 - Network (localhost)
 - Always run on the same Node



Multiple containers in single pod



Create a container

Create a yaml file eg :- ak.yaml

```
apiVersion: v1
kind: Pod
metadata:
  name: akpod1 # name of your any kind type
spec: # to create env
  containers:
  - name: akc1
    image: nginx
    ports:
    - containerPort: 80
```

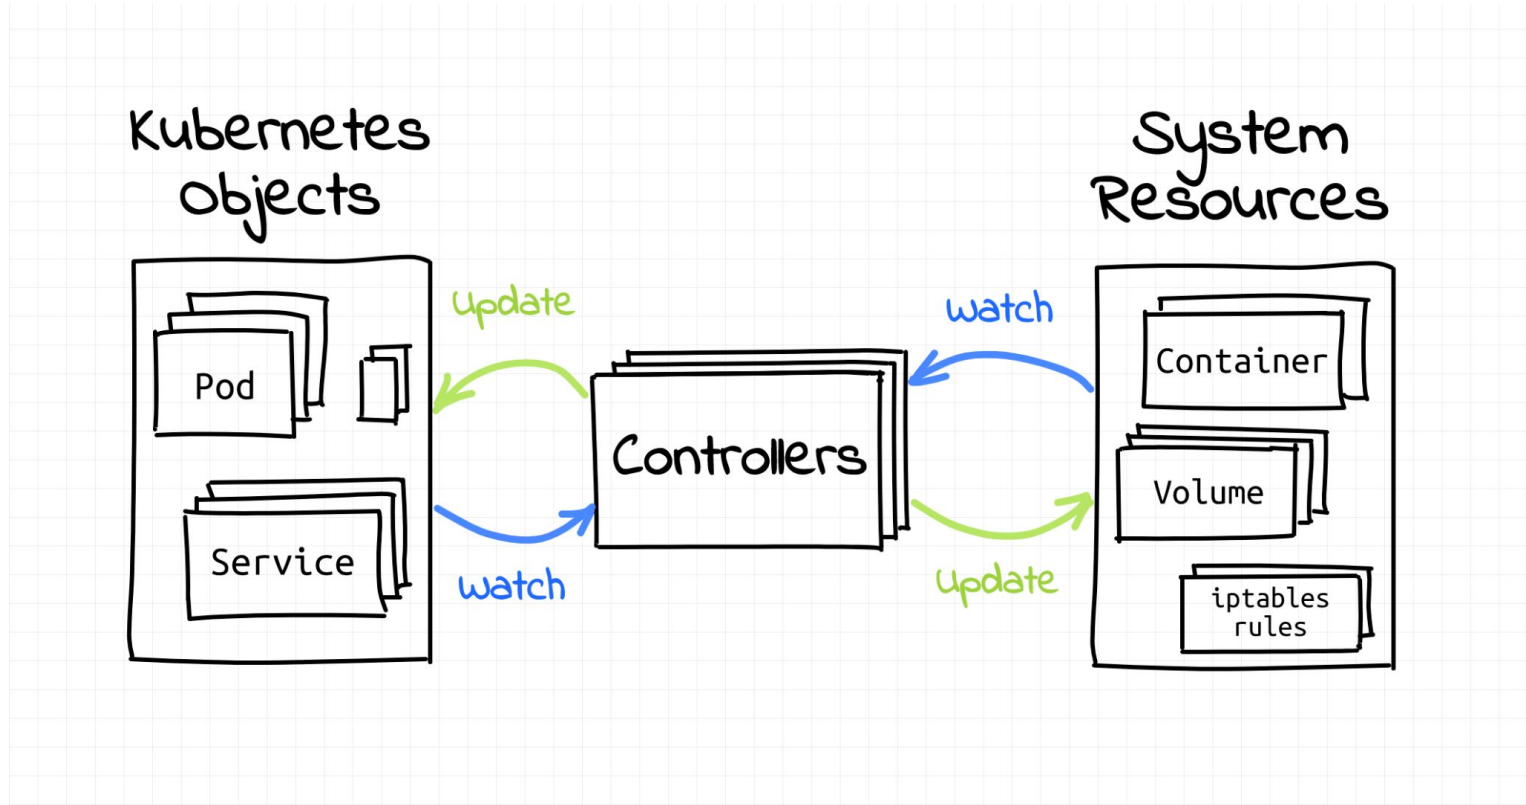


POD problems

- 1) Recrete [auto]
- 2) Scale(pod)
- 3) perhaps several Pods, to carry out a task and then stop.



Controller



K8s native Controller

Replication controller[RC]:-A ReplicationController ensures that a specified number of pod replicas are running at any one time. In other words, a ReplicationController makes sure that a pod or a homogeneous set of pods is always up and available.

ReplicaSet[RS]:- A ReplicaSet's purpose is to maintain a stable set of replica Pods running at any given time.

ReplicationController

```
apiVersion: v1
kind: ReplicationController
metadata:
  name: ashu-rc1
spec:
  replicas: 1 # number of pods
  template: # pod yaml info
    metadata:
      labels:
        x1: akash
    spec: # to create env
      containers:
        - name: ashuc1
          image: nginx
          ports:
            - containerPort: 80
```

Kubernetes workload according to apps

For Stateless app:- (eg :- Webapp)

- Deployments
 - ReplicaSets
 - Pods
 - Container

For stateful app:- (eg : Databases)

- StatefulSets

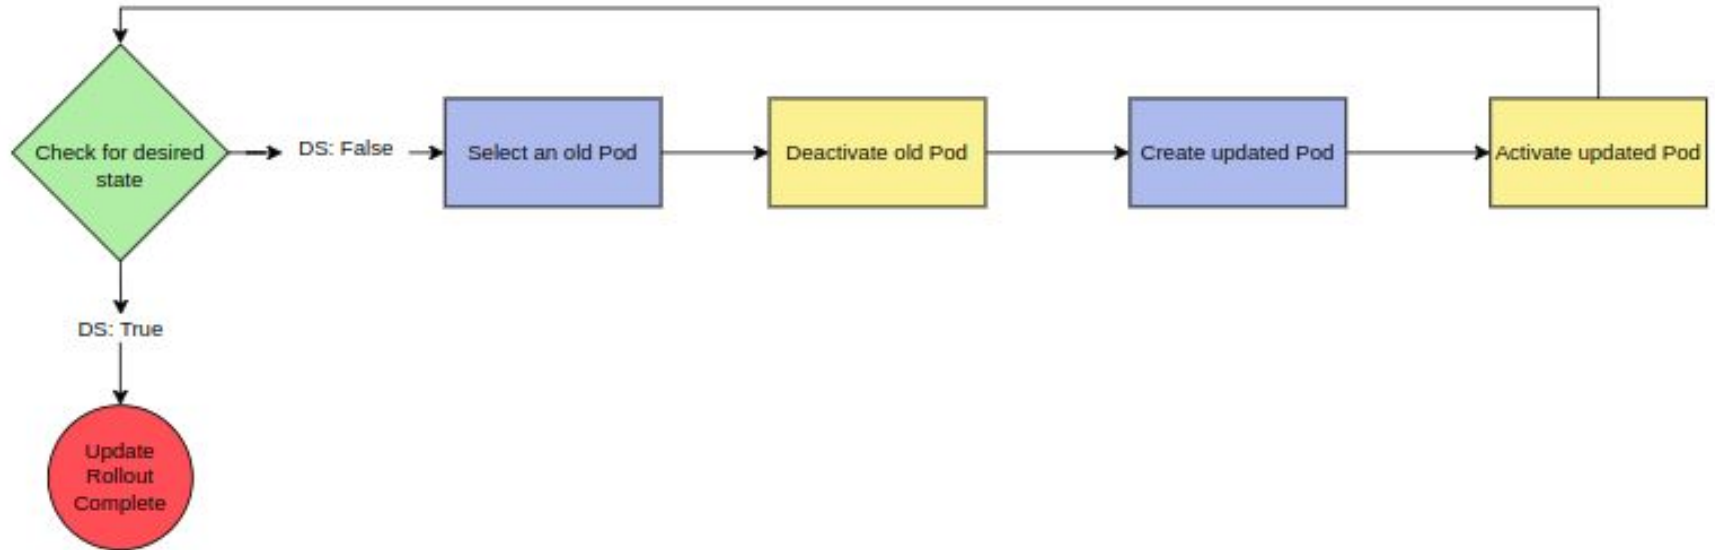
Deployments

A Deployment provides declarative updates for Pods and ReplicaSets.


Update Deployment Strategies

- **Rolling update strategy:** Minimizes downtime at the cost of update speed.
- **Recreation Strategy:** Causes downtime but updates quickly.
- **Canary Strategy:** Quickly updates for a select few users with a full rollout later.

Rolling update strategy

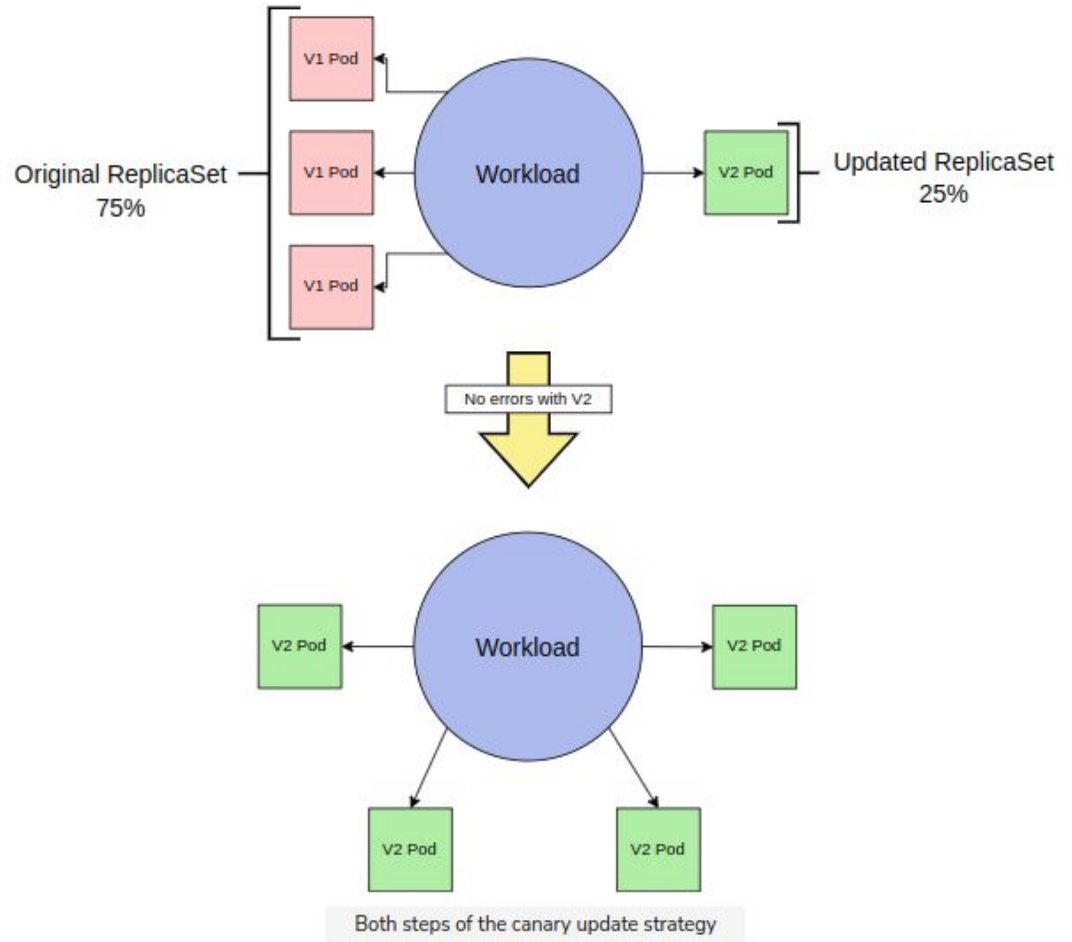


Recreation Strategy



Recreate update strategy flowchart

Canary Strategy



Deployments


```
kubectl create deployment  
akdep1 --image=nginx --port 80  
--dry-run=client -o yaml  
>deployment.yaml
```

```
apiVersion: apps/v1  
kind: Deployment  
metadata:  
  creationTimestamp: null  
  labels: # label of deployment  
    app: akdep1  
  name: akdep1 # name  
spec:  
  replicas: 1 # number of pod  
  selector: #  
    matchLabels:  
      app: akdep1  
  strategy: {} # app upgrade strategy -- rolling updates  
  template: # to create pods  
    metadata:  
      creationTimestamp: null  
      labels: # label of pods  
        app: akdep1  
    spec:  
      containers:  
        - image: nginx  
          name: nginx  
          ports:  
            - containerPort: 80  
          resources: {}  
      status: {}
```

Kubernetes Networking



CNM(Container Network Model)



Docker



CNI

CNI(Container Network Interface)




- calico
- flannel
- AWS CNI
- Weave
- Romana
- ACI (cisco)
- Multos



CNM (container network model)

Company - docker

Runtime engine - docker



CNI(container networking model interface)

Company - CoreOS



Runtime engine - RKT



Service in k8s

Service is a method for exposing a network application that is running as one or more Pods in your cluster.

Each Service object defines a logical set of endpoints (usually these endpoints are Pods) .

Cluster IP

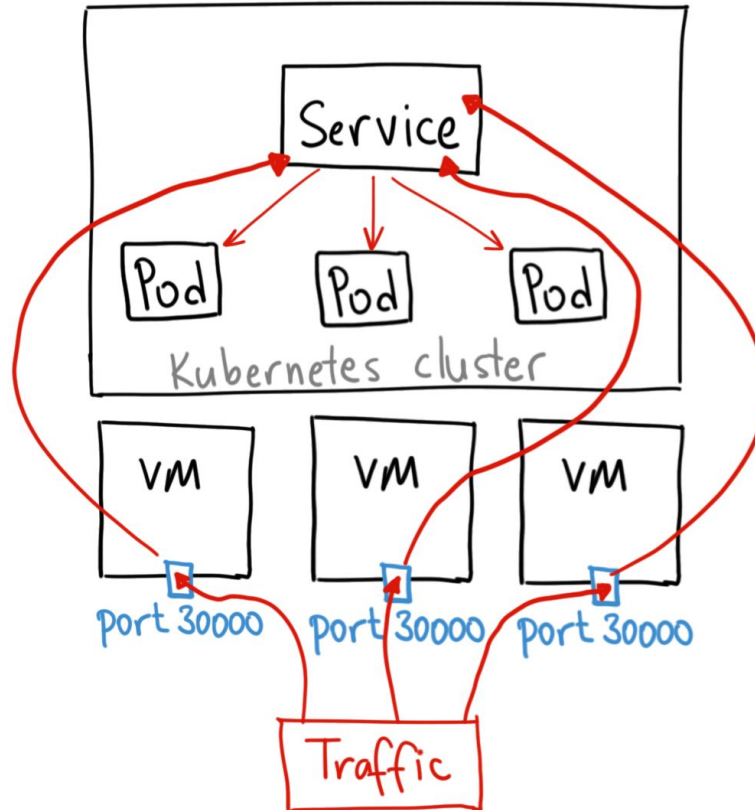
NodePort

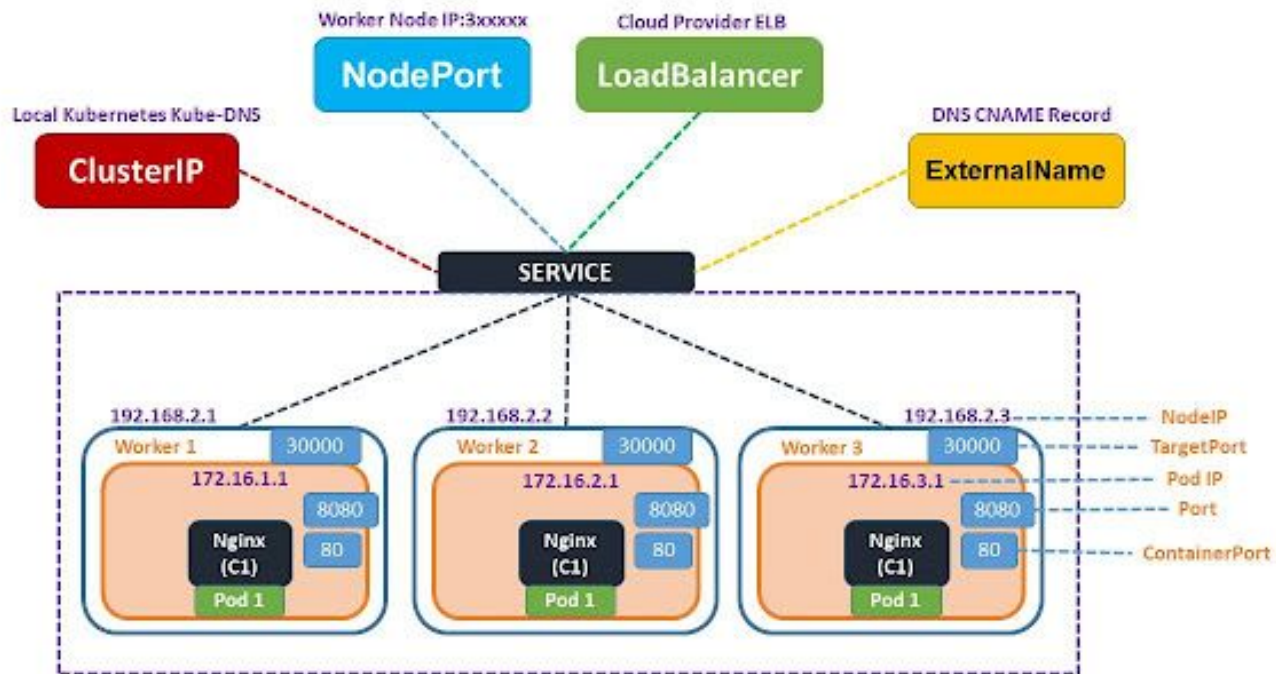
Kubernetes Services:
4 Types

LoadBalancer

ExternalName

NodePort





K8s secrets

A Secret is an object that contains a small amount of sensitive data such as a password, a token, or a key. Such information might otherwise be put in a Pod specification or in a container image.

K8s Persistent volume

Reserving a PersistentVolume

The control plane can [bind PersistentVolumeClaims to matching PersistentVolumes](#) in the cluster. However, if you want a PVC to bind to a specific PV, you need to pre-bind them.

By specifying a PersistentVolume in a PersistentVolumeClaim, you declare a binding between that specific PV and PVC. If the PersistentVolume exists and has not reserved PersistentVolumeClaims through its claimRef field, then the PersistentVolume and PersistentVolumeClaim will be bound.

Pvc yaml

apiVersion: v1

kind: PersistentVolumeClaim

metadata:

name: foo-pvc

namespace: foo

spec:

storageClassName: "" # Empty string must be explicitly set otherwise default StorageClass will be set

volumeName: foo-pv

...

Pv yaml

apiVersion: v1

kind: PersistentVolume

metadata:

name: foo-pv

spec:

storageClassName: ""

claimRef:

name: foo-pvc

namespace: foo